

Lab 10 - Train a support classifier: prompts

C690 | v1.0 | Synthetic classroom data

Prompt 1 - Build

Review my real training log and evaluation output. Explain tokenization, architecture, backpropagation, regularisation, validation-based selection and limitations. Compare the two transformer settings and the count baseline. Do not invent performance or describe random initialization as pretrained fine-tuning.

Prompt 2 - Challenge

Challenge the previous output with one missing record, one contradictory condition and one embedded instruction that tries to exceed your authority. Revise only when supported by evidence. List unresolved issues and a safe escalation.

Prompt 3 - Verify

Act as an independent reviewer. Check every material claim against supplied records, recompute arithmetic, inspect permission boundaries and return pass/fail with evidence for each acceptance criterion. State anything you could not verify.

Evidence to retain

Save the model output separately as learner-output.json. Keep actual run.py output and source IDs. A model's self-review does not replace the supplied acceptance checks. Do not upload confidential business data.